

# A/B Testing Project

Dataset - [https://www.kaggle.com/datasets/zhangluyuan/ab-testing/code?select=ab\\_data.csv](https://www.kaggle.com/datasets/zhangluyuan/ab-testing/code?select=ab_data.csv)

Reference Course - <https://www.udacity.com/course/ab-testing--ud257>

Disclaimer - This is just to showcase some skill, the dataset needs more features to make better explainable conclusions.

## 1. Experiment Introduction and Hypothesis Testing

The aim is to formulate and test the hypothesis for comparing the conversion rate between Old web page vs New web page.

Two Tailed test:

$$H_0: p = p_0$$

$$H_a: p \neq p_0$$

At confidence interval of 95% -  $\alpha = 0.05$

## 2. Data and Preprocessing

```
import pandas as pd
import numpy as np
```

Import Data

```
df = pd.read_csv('ab_data.csv')
```

Independent Variables: -

Control Group is shown Old web page.

Treatment Group is shown New web page.

Dependent Variable: -

Converted (0 or 1) represents a buy or click.

```
df.head()
```

|           | user_id | timestamp                  | group   | landing_page |
|-----------|---------|----------------------------|---------|--------------|
| converted |         |                            |         |              |
| 0         | 851104  | 2017-01-21 22:11:48.556739 | control | old_page     |
| 0         |         |                            |         |              |
| 1         | 804228  | 2017-01-12 08:01:45.159739 | control | old_page     |

```

0
2   661590   2017-01-11 16:55:06.154213   treatment   new_page
0
3   853541   2017-01-08 18:28:03.143765   treatment   new_page
0
4   864975   2017-01-21 01:52:26.210827   control     old_page
1

df.shape
(294478, 5)

```

Check for null values and data type

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 294478 entries, 0 to 294477
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id         294478 non-null   int64
1   timestamp       294478 non-null   object
2   group           294478 non-null   object
3   landing_page    294478 non-null   object
4   converted       294478 non-null   int64
dtypes: int64(2), object(3)
memory usage: 11.2+ MB

```

Data Description

```

df.describe()

          user_id      converted
count  294478.000000  294478.000000
mean    787974.124733      0.119659
std      91210.823776      0.324563
min      630000.000000      0.000000
25%     709032.250000      0.000000
50%     787933.500000      0.000000
75%     866911.750000      0.000000
max     945999.000000      1.000000

pd.crosstab(df['group'], df['landing_page'])

landing_page  new_page  old_page
group
control        1928    145274
treatment     145311      1965

df.converted.mean()

```

0.11965919355605512

Check for duplicate user\_id

```
session_counts = df['user_id'].value_counts(ascending=False)
multi_users = session_counts[session_counts > 1].count()

print(f'There are {multi_users} users that appear multiple times in
the dataset')
```

There are 3894 users that appear multiple times in the dataset

Remove duplicate user\_id

```
users_to_drop = session_counts[session_counts > 1].index

df = df[~df['user_id'].isin(users_to_drop)]
print(f'The updated dataset now has {df.shape[0]} entries')
```

The updated dataset now has 286690 entries

Re-check for duplicate user\_id

```
df[df.user_id.duplicated(keep=False)]
```

Empty DataFrame

Columns: [user\_id, timestamp, group, landing\_page, converted]

Index: []

Assumptions to calculate sample size: -

Lets Assume the baseline conversion rate is 11% based on the past data and through the experiment the organization is aiming for 14% conversation rate.

The function `statsmodels.stats.proportion_effectsize(0.11, 0.14)` is calculating the effect size based on a proportion for two independent samples, in this case with proportions of 0.11 and 0.14.

The function `statsmodels.stats.NormalIndPower().solve_power(effect_size, power=0.8, alpha=0.05, ratio=1)` is used to calculate the sample size required for the study. Here `power=0.8` and `alpha=0.05` are the statistical power and significance level used for the calculation, respectively, and `ratio=1` indicates that the sample sizes for the two groups are equal.

In the context of hypothesis testing in statistics, power is the probability that a test correctly rejects the null hypothesis when the alternative hypothesis is true. The power of a statistical test is generally defined as  $1 - \beta$ , where  $\beta$  is the probability of Type II error.

Type II error occurs when the null hypothesis is false, but we fail to reject it. So, the power of a test is essentially the test's ability to detect an effect if there is one.

Finally, `required_n = ceil(required_n)` rounds up the number to ensure it is a whole number, since we can't have a fraction of a sample.

```
import statsmodels.stats.api as sms
from math import ceil

effect_size = sms.proportion_effectsize(0.11, 0.14)    # Calculating
effect size based on our expected rates

required_n = sms.NormalIndPower().solve_power(
    effect_size,
    power=0.8,
    alpha=0.05,
    ratio=1
)                                                       # Calculating
sample size needed

required_n = ceil(required_n)                          # Rounding up
to next whole number

print(required_n)

1902
```

### Sample Collection from both Control Group and Treatment Group

```
control_sample = df[df['group'] == 'control'].sample(n=required_n,
random_state=22)
treatment_sample = df[df['group'] == 'treatment'].sample(n=required_n,
random_state=22)

ab_test = pd.concat([control_sample, treatment_sample], axis=0)
ab_test.reset_index(drop=True, inplace=True)

ab_test.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3804 entries, 0 to 3803
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   user_id         3804 non-null   int64
1   timestamp       3804 non-null   object
2   group           3804 non-null   object
3   landing_page    3804 non-null   object
4   converted       3804 non-null   int64
dtypes: int64(2), object(3)
memory usage: 148.7+ KB

ab_test['group'].value_counts()

control      1902
treatment    1902
Name: group, dtype: int64
```

### Mean and Standard Deviation of Control Group and Treatment Group

```
import scipy.stats as stats

conversion_rates = ab_test.groupby('group')['converted']

std_p = lambda x: np.std(x, ddof=0)           # Std. deviation of
the proportion
se_p = lambda x: stats.sem(x, ddof=0)         # Std. error of the
proportion (std / sqrt(n))

conversion_rates = conversion_rates.agg([np.mean, std_p, se_p])
conversion_rates.columns = ['conversion_rate', 'std_deviation',
'std_error']

conversion_rates.style.format('{:.3f}')

<pandas.io.formats.style.Styler at 0x25c4c2005b0>
```

The Mean and Standard Deviation of Control Group and Treatment Group are fairly close indicating not much of significance difference between groups even before test.

### 3. Plot of Conversion for Control Group and Treatment Group

```
import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,6))

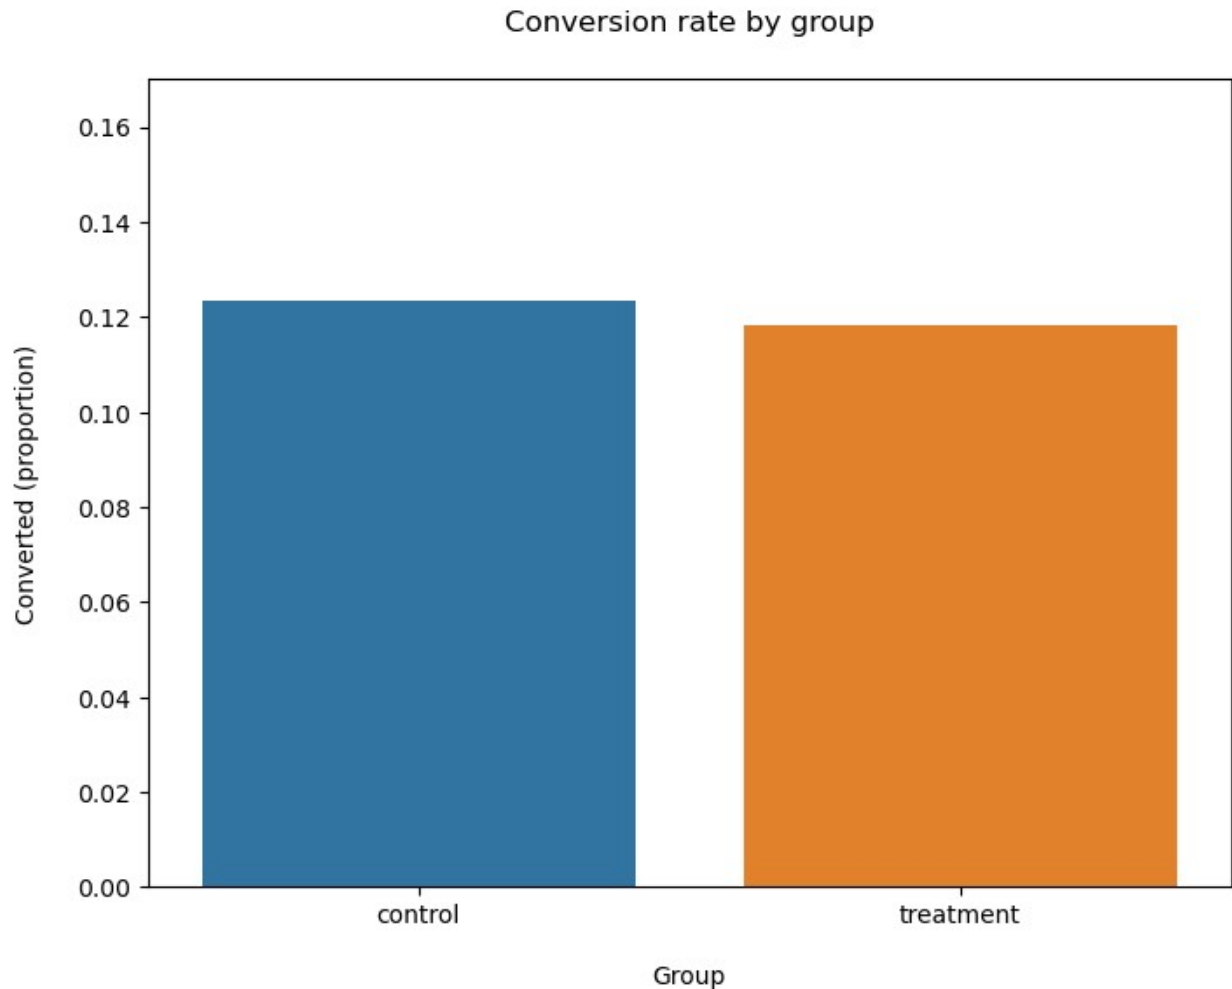
sns.barplot(x=ab_test['group'], y=ab_test['converted'], ci=False)

plt.ylim(0, 0.17)
plt.title('Conversion rate by group', pad=20)
plt.xlabel('Group', labelpad=15)
plt.ylabel('Converted (proportion)', labelpad=15);
```

C:\Users\achal\AppData\Local\Temp\ipykernel\_10568\3362076549.py:6:  
FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=('ci', False)` for the same effect.

```
sns.barplot(x=ab_test['group'], y=ab_test['converted'], ci=False)
```



#### 4. Testing the Hypothesis: -

The goal is to determine if the proportion of successes (conversions) is significantly different between the two groups. The test statistic is a Z-score (Z-stat).

The two-proportion Z-test is used when you have two samples and you want to see if the success rates are significantly different from each other.

At other times, a binomial test may be used when you have one sample and you want to compare the observed success rate to a hypothesized success rate.

```
from statsmodels.stats.proportion import proportions_ztest,
proportion_confint

control_results = ab_test[ab_test['group'] == 'control']['converted']
treatment_results = ab_test[ab_test['group'] == 'treatment']
['converted']

n_con = control_results.count()
```

```

n_treat = treatment_results.count()
successes = [control_results.sum(), treatment_results.sum()]
nobs = [n_con, n_treat]

z_stat, pval = proportions_ztest(successes, nobs=nobs)
(lower_con, lower_treat), (upper_con, upper_treat) =
proportion_confint(successes, nobs=nobs, alpha=0.05)

print(f'z statistic: {z_stat:.2f}')
print(f'p-value: {pval:.3f}')
print(f'ci 95% for control group: [{lower_con:.3f}, {upper_con:.3f}]')
print(f'ci 95% for treatment group: [{lower_treat:.3f},
{upper_treat:.3f}]')

z statistic: 0.50
p-value: 0.619
ci 95% for control group: [0.109, 0.138]
ci 95% for treatment group: [0.104, 0.133]

```

## 5. Conclusions: -

For a two-tailed test at a 95% confidence level (significance level of  $\alpha = 0.05$ ), the critical Z values are approximately -1.96 and +1.96. suggesting that the result is not statistically significant at the 95% confidence level and fails to reject the null hypothesis.

The p-value of  $0.619 > 0.05$  and thus, insignificant to reject null hypothesis  $H_0$ .

The confidence values of treatment group are not approaching to 14% aim expectation.

The confidence values of treatment group does not show increment compared to control group instead has only slightly decreased.

## Thank you